

# Temporal Data II (Pandas + Analysis)

## SI370

March 16, 2017

Eytan Adar

# Last Time: Temporal data I

- What is temporal data?
- How do we visualize it?
  - Time series
  - Clustering
  - Event-driven analysis
- Manipulation
  - Pandas examples

# Case 7

- Abstract task:
  - Find repeating patterns (seasonality) in time series datasets
  - Isolate the periodicity
  - Compare aligned time series

Spiral visualizations  
Carlis et al., 1998




Figure 2. An indented spiral, with spokes, showing monthly consumption percentages for *Baphia Capparidifolia* during the period 1980 – 1988.





Figure 7. Tightening a spiral view of sound data from five instruments. From left to right the structure of the sound reveals itself.

Spiral visualizations  
Carlis et al., 1998



Spiral

Bars  Period: 3

Births

Downloads

# Summary

- Identify the different domain goals
- Isolate the tasks (there are many)
  - Select encodings that support the tasks
  - Standard types are a good place to start
    - Additional retinal variables often useful
  - Some less standard types might also be appropriate
- Now for some Pandas...





# Today: Temporal data

- What is temporal data?
- How do we visualize it?
  - Time series
  - Clustering
  - Event-driven analysis
- Manipulation
  - Pandas examples
- Analysis
  - Trends and seasonality
  - Cross-correlation
  - Autocorrelation



# A few key concepts

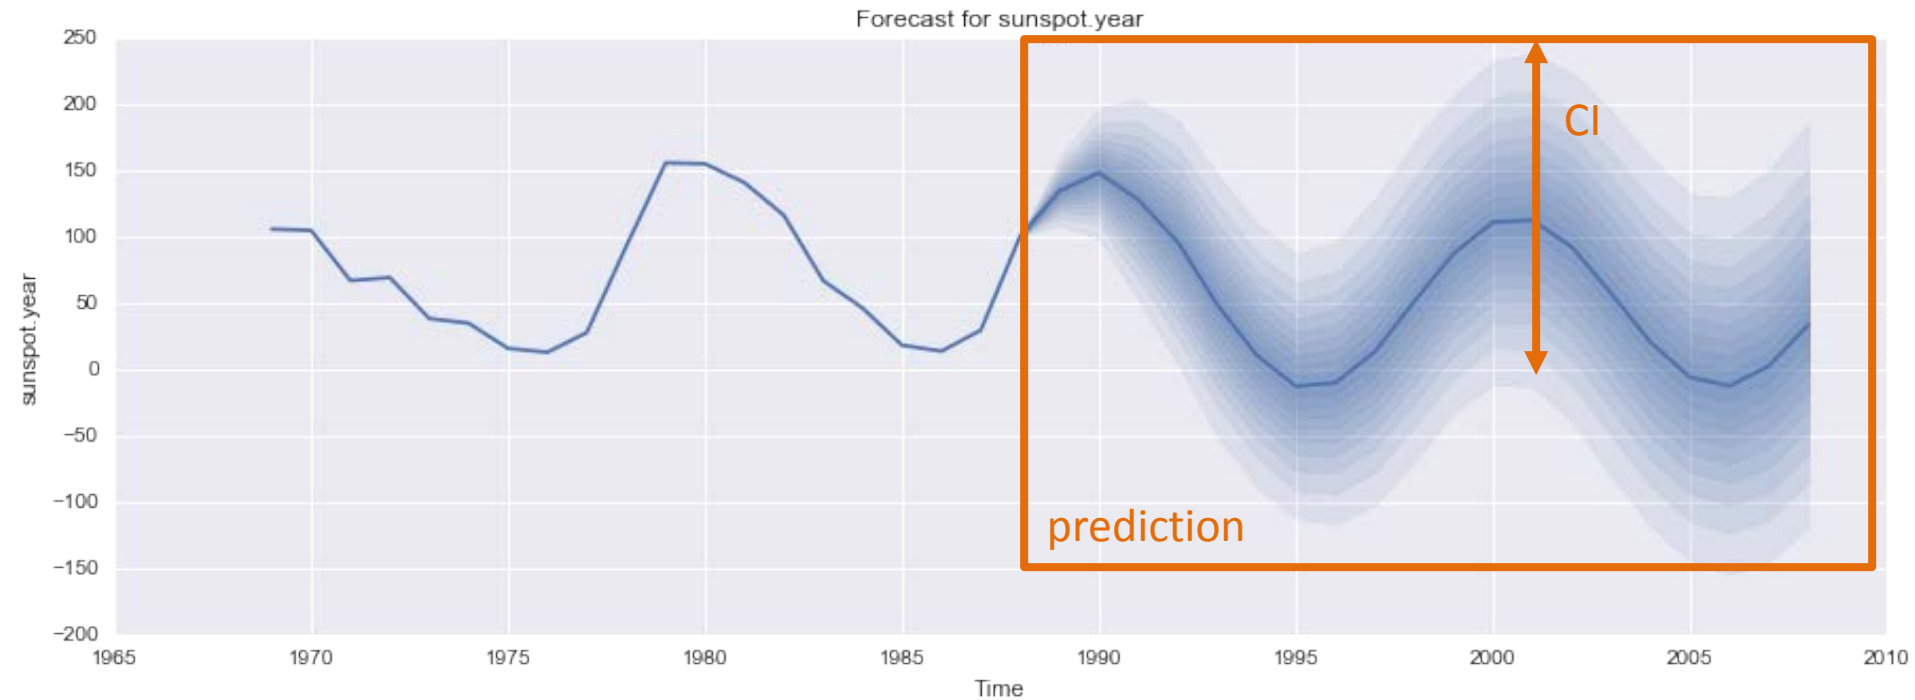
- “Univariate” time-series
  - We care about how it breaks down



# A few key concepts

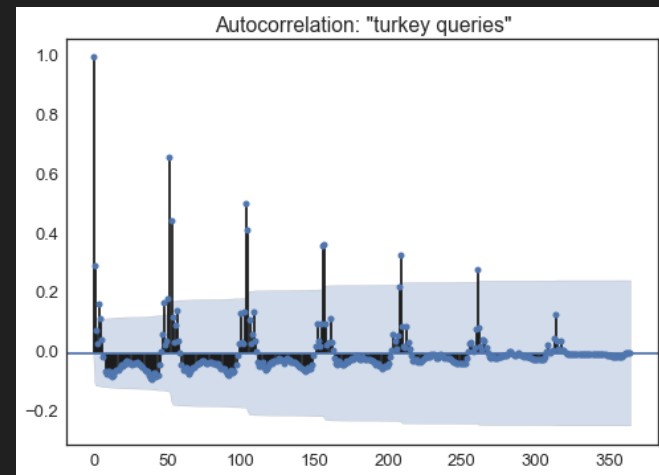

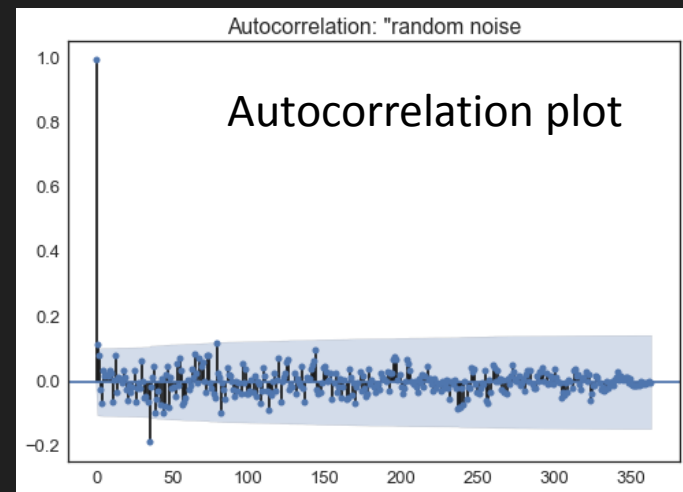

- “Univariate” time-series
  - We care about how it breaks down
    - Seasonality, trend, residuals
  - We care about what value it will take in the future

## ARIMA model for sunspots



# A few key concepts

- “Univariate” time-series
  - We care about how it breaks down
    - Seasonality, trend, residuals
  - We care about what value it will take in the future
    - For example: ARIMA
      - Determine parameters (ACF, PACF)



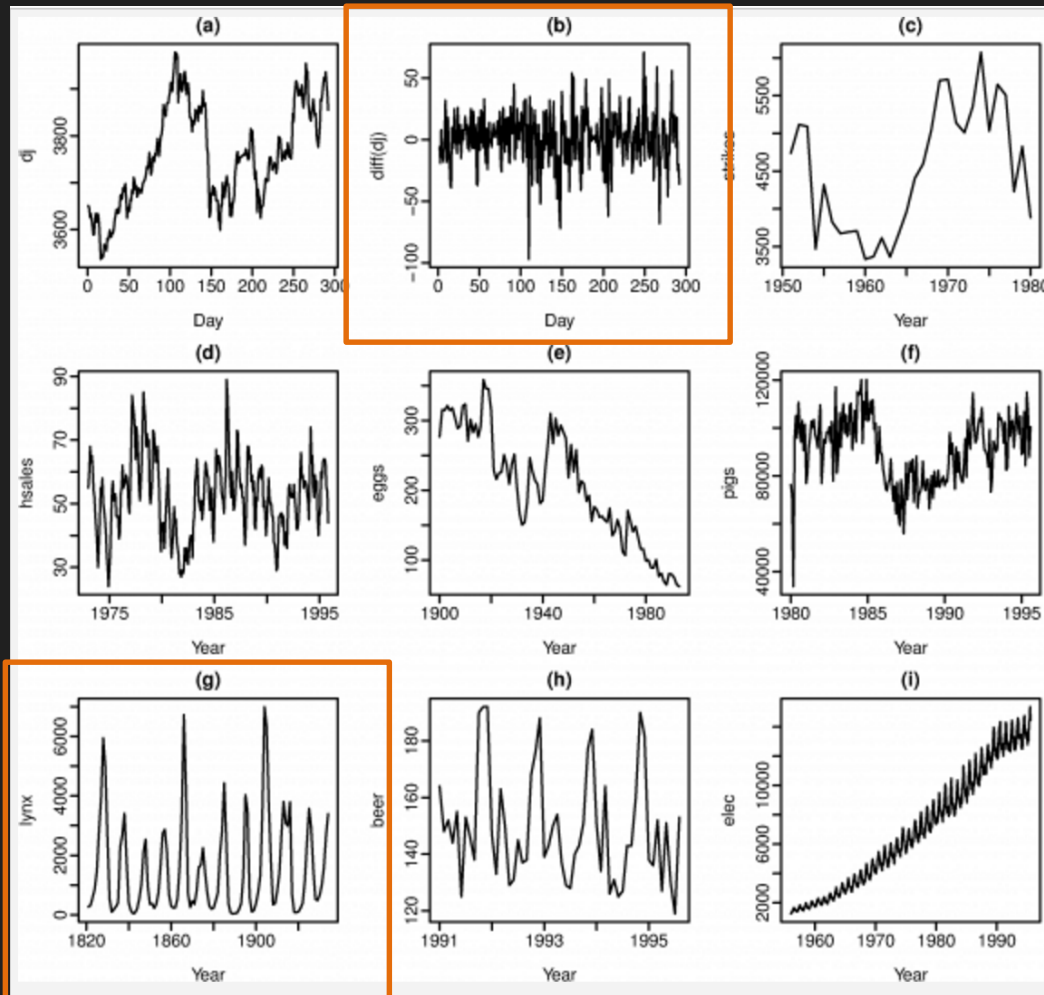
# Auto-correlation

- Is the data noise?
- Compare the data to itself at different lags
  - If we only see one “peak” (basically at lag 0)
    - Random
  - If we see other peaks/hills
    - Indication of seasonality
    - Or that value at  $t_n$  “depends on value”  $t_{n-1}$

# A few key concepts

- “Univariate” time-series
  - We care about how it breaks down
    - Seasonality, trend, residuals
  - We care about what value it will take in the future
    - For example: ARIMA
    - Some assumptions about this: stationary





**(a)** Dow Jones index on 292 consecutive days; **(b)** Daily change in Dow Jones index on 292 consecutive days; **(c)** Annual number of strikes in the US; **(d)** Monthly sales of new one-family houses sold in the US; **(e)** Price of a dozen eggs in the US (constant dollars); **(f)** Monthly total of pigs slaughtered in Victoria, Australia; **(g)** Annual total of lynx trapped in the McKenzie River district of northwest Canada; **(h)** Monthly Australian beer production; **(i)** Monthly Australian electricity production

# A few key concepts


- “Univariate” time-series
  - We care about how it breaks down
    - Seasonality, trend, residuals
  - We care about what value it will take in the future
    - For example: ARIMA
    - Some assumptions about this: stationary
      - We can make data stationary

# A few key concepts

- Bi-variate time-series analysis
  - How related are two time series?
    - Cross-correlation
      - A bit like autocorrelation

# Cross-Correlation



- Works much the same as standard correlation.
  - $\text{corr}(x,y) \rightarrow$  treats each pair as “point”
  - $t1 \rightarrow x1,y1; t2 \rightarrow x2,y2 ; \dots$
- One exception:
  - Lags/delays
  - Often one series is in advance of other



---

# Cross-Correlation

- Works much the same as standard correlation.
  - $\text{corr}(x,y) \rightarrow$  treats each pair as “point”
  - $t1 \rightarrow x1,y1; t2 \rightarrow x2,y2 ; \dots$
- One exception:
  - Lags/delays
  - Often one series is in advance of other
- We can include a “lag”
  - Often part of utility function... provide min/max lag
  - Find max value (above some critical point)



# Summary

- Temporal data sets are ubiquitous
- Need to carefully consider the problem
  - Different ways to structure the data
  - Different ways to do EDA
    - Visual forms and statistical
  - Statistics are wonky... data needs to be “cleaned” differently than other datasets
- Now... more Pandas